

Natural Language Processing

Lecture 01: Introduction to NLP

Valentin Malykh, MTS AI



Autumn 2024

The course is delivered at ITMO University, Saint-Petersburg &
Bauman University, Moscow

Content

- 1 About the course
- 2 Research questions and NLP tasks
- 3 Grammars and Automata
- 4 Text segmentation and morphology analysis
- 5 Word frequency and collocations

Content

- 1 About the course
- 2 Research questions and NLP tasks
- 3 Grammars and Automata
- 4 Text segmentation and morphology analysis
- 5 Word frequency and collocations

Acknowledgements

- Dr. Qun Liu, Huawei Noah's Ark lab
- Dr. Constantin Korikov, Huawei Noah's Ark lab
- Tasnima Sadekova, Huawei Noah's Ark lab
- Salavat Garifullin, ODS
- students from previous runs

Logistics

- Instructor: Dr. Valentin Malykh
- TAs: Salavat Garifullin for ODS.ai & others
- Lecture Time: 19.00, Thursdays
- Seminar Time: 19.00, Tuesdays
- Location: Online
- Slides: will be available at the course platform before each class.

Grading policy

- Quizzes – 8 x 4
- Assignments – 25 x 2
- Final project – 60

Course description

Natural Language Processing (NLP) is a domain of research whose objective is to analyze and understand human languages and develop technologies to enable human machine interactions with natural languages. NLP is an interdisciplinary field involving linguistics, computer sciences and artificial intelligence. The goal of this course is to provide students with comprehensive knowledge of NLP. Students will be equipped with the principles and theories of NLP, as well as various NLP technologies, including rule-based, statistical and neural network ones. After this course, students will be able to conduct NLP research and develop state-of-the-art NLP systems.

Content

- 1 About the course
- 2 Research questions and NLP tasks**
- 3 Grammars and Automata
- 4 Text segmentation and morphology analysis
- 5 Word frequency and collocations

Natural language processing in Wikipedia

Natural language processing (NLP) is a subfield of linguistics, computer science, information engineering, and artificial intelligence concerned with the interactions between computers and human (natural) languages, in particular how to program computers to process and analyze large amounts of natural language data.

Synonyms of NLP

- **Computational Linguistics**
- **Natural Language Processing**
- **Natural Language Understanding**
- **Human Language Processing**

- Subtleties

Computational Linguistics is more regarded as a branch of Linguistics, whose main purpose is to understand the mechanism of human languages by means of computing

Synonyms of NLP

- **Computational Linguistics**
- **Natural Language Processing**
- **Natural Language Understanding**
- **Human Language Processing**

- Subtleties

Natural Language Processing is a branch of computer sciences and artificial intelligent, whose main purpose is to develop technologies to enable human-computer interactions using human languages

Synonyms of NLP

- **Computational Linguistics**
- **Natural Language Processing**
- **Natural Language Understanding**
- **Human Language Processing**

- **Subtleties**

Natural Language Understanding is one of the two main challenges in Natural Language Processing, while the other is **Natural Language Generation**.


Synonyms of NLP

- **Computational Linguistics**
- **Natural Language Processing**
- **Natural Language Understanding**
- **Human Language Processing**

- Subtleties

Human Language Technologies mainly refer to NLP technologies, but may also include other language related technologies, include speech technologies, optical character recognition (OCR), computer typesetting, etc.

NLP as an interdisciplinary study




Understanding human languages is not easy

- We are getting used to the fact that human beings can understand each other using language communication.
- Although it is a natural result of evolution for human to obtain the language competence.
- It seems to be a miracle, due to its complexity.
- No other species in this planet can use languages at the degree as humans do.
- The mechanism behind human languages is not fully discovered.
- Understanding human languages by computer is difficult.

Understanding human languages is not easy

Tell my wife I love her!



**From Husband:
I love her!**



Research questions

- How humans understand each other by using language communication?
- Is it possible to simulate human language behaviors without understanding language mechanisms?

The way of NLP research

- Unlike linguists who develop numerous theories to explain the language mechanisms, NLP researchers try to simulate human language behaviors by computing, not necessary to understand the language mechanisms.


A brief history of NLP

- 1960s-1990s: Rule-based approaches
- 1990s-2010s: Statistical approaches
- 2010s-present: Neural network (deep learning) approaches

Holy grails of NLP

- Accurate machine translation between human languages
- Free conversation between humans and computers

Accurate machine translation




The Tower of Babel




Oil painting by Pieter Bruegel the Elder, 1563, from Wikipedia

Free human machine conversation




Turing test




By Juan Alberto Sánchez Margallo, CC BY 2.5, from Wikipedia

Languages, human minds and the world



NLP Tasks

1 Bit Bits to Character Encoding															 www.innerdoc.com		75 App Interactive App Creation					
2 Typ Manual Typewriting	8 Man Manual Annotation					29 Pri Price Parser											63 Nex Next Token Prediction	69 Rel Relation Extraction	76 Ann Annotated Text Visualization			
3 Str Loading a Structured Datasets	9 Act Annotation with Active Learning	14 Tok Tokenization	19 Ste Stemming	24 Ngr N-grams	30 Geo Geocoding											43 Trn Training Models	48 Spa Spam Detection	53 Key Keyword Extraction	58 Syn Wordnet Synsets	64 Rep Report Writing	70 Qan Question Answering	77 Wcl Wordcloud
4 Cor Generating a Corpus	10 Pro Training Data Provider	15 Voc Vocabulary Building	20 Lem Lemmatization	25 Phr Rulebased Phrasematcher	31 Tmp Temporal Parser	35 Sen Sentencizer	39 Ded Deduplication	44 Tst Evaluating Models	49 Sed Sentiment and Emotion Detection	54 Esu Extractive Summarization	59 Dst Distance Measures	65 Tra Machine Translation	71 Cha Chatbot Dialogue	78 Emb Word Embedding Visualization								
5 Api Loading from API	11 Cro Crowdsourcing Marketplace	16 Mor Morphological Tagger	21 Nrm Normalization	26 Chu Dependency Nounchunks	32 Nel Named Entity Linking	36 Par Paragraph Segmentation	40 Raw Raw Text Cleaning	45 Exp Explaining Models	50 Int Intent Classification	55 Top Topic Modeling	60 Sim Document Similarity	66 Asu Abstractive Summarization	72 Sem Semantic Search Indexing	79 Tim Events on Timeline								
6 Scr Text and File Scraping	12 Aug Textual Data Augmentation	17 Pos Part-of-Speech Tagger	22 Spl Spell Checker	27 Ner Named Entity Recognition	33 Crf Conference Resolution	37 Grm Grammar Checker	41 Met Meta-Info Extractor	46 Dpl Deploying Models	51 Cls Text Classification	56 Tre Trend Detection	61 Dis Distributed Word Representations	67 Prp Paraphrasing	73 Kno Knowledge Base Population	80 Map Locations on Geomap								
7 Ext Text Extraction and OCR	13 Rul Rulebased Training Data	18 Dep Dependency Parser	23 Neg Negation Recognizer	28 Abr Abbreviation Finder	34 Anm Text Anonymizer	38 Rea Readability Scoring	42 Lng Language Identification	47 Mon Monitoring Models	52 Mlc Multi-Label Multi-Class Classification	57 Out Outlier Detection	62 Con Contextualized Word Representations	68 Lon Long Text Generation	74 Edi E-Discovery and Media Monitoring	81 Gra Knowledge Graph Visualization								
Source Data Loading	Training Data Generation	Word Parsing	Word Processing	Phrases and Entities	Entity Enriching	Sentences and Paragraphs	Documents	Model Development	Supervised Classification	Unsupervised Signaling	Similarity	Natural Language Generation	Systems	Information Visualization								



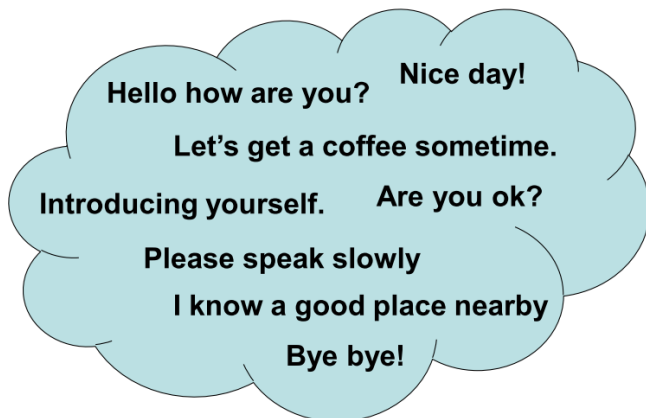
www.innerdoc.com

	Classification	Tagging	Generation
n-grams	TF-IDF	regex	templates
word embeddings	word2vec	word2vec	
CNN	CNN	CNN	ConvSeq2Seq
RNN	LSTM	LTSM	LSTM
Transformers	BERT	BERT	T5
LLM	LLaMa	LLaMa	LLaMa

Content

- 1 About the course
- 2 Research questions and NLP tasks
- 3 Grammars and Automata**
- 4 Text segmentation and morphology analysis
- 5 Word frequency and collocations

How can we define a language?



How can we define a language?

- A language can be defined as the set of sentences which can be accepted by the speakers of that language.
- It is not possible to define a natural language by enumerate all the sentences, because the number of sentences in a natural languages is infinite.
- Two feasible ways to define a language with infinite sentences:
 - By a Grammar
 - By an Automaton




Define a language by an automaton (1)

- An automaton A is a abstract machine which:
 - Takes a symbol sequence S as input, and determines if A will *accept* or *reject* S .
 - Has a finite number of states and a finite number of actions.
 - At each time step, S is in a state, and points to a position in S .
 - The current state and current symbol determines the action which A will execute, which determines the next state of A and the next position of S where A will point to.
 - Given a input S , A will run until it stops, and the final state of A determines if A will *accept* or *reject* S .

Define a language by an automaton (2)

- A language L can be defined by an automaton A as:
 - A word sequence S is a sentence of L , if and only if: when we input S to A , A will stop in a finite number of time steps at an *accept* state.

Turing machine



Turing machine

A Turing machine consists of: (to be continued)

- A *tape* divided into cells, one next to the other. Each cell contains a symbol from some finite alphabet. The alphabet contains a special blank symbol and some other symbols. The *tape* is assumed to be arbitrarily extendable to the left and to the right.
- A *read/write head* that can read and write symbols on the *tape* and move the *tape* left and right one (and only one) cell at a time.

Turing machine



- A Turing machine consists of: (continued)
 - A *state register* that stores the state of the Turing machine, one of finitely many. Among these is the special *start state* with which the state register is initialized.
 - A finite *table* of instructions that, given the *state* the machine is currently in and the symbol it is reading on the *tape*, tells the machine to do the following in sequence:
 - Either erase or write a symbol.
 - Move the *head* to the left or right cell.
 - Assume the same or a *new state* as prescribed.

Linear bounded automaton

A linear bounded automaton is a Turing Machine that satisfies the following three conditions:

- Its input alphabet includes two special symbols, serving as *left and right endmarkers*.
- Its *transitions* may not print other symbols over the *endmarkers*.
- Its *transitions* may neither move to the left of the *left endmarker* nor to the right of the *right endmarker*.

Finite state automaton / machine (FSA/FSM)



Finite state automaton / machine (FSA/FSM)

A Finite State Automaton (FSA), or Finite State Machine (FSM), consists of:

- A finite number of *states*, while the FSM can be in one *states* at each given time;
- A *head* which read a symbol from a sequence of symbols as the *input*. The *head* always goes to the next symbol at the next time step;
- A *transition* matrix which determines the next *states* of the FSM according to the current *states* and the current symbol.

Content

- 1 About the course
- 2 Research questions and NLP tasks
- 3 Grammars and Automata
- 4 Text segmentation and morphology analysis**
- 5 Word frequency and collocations

Text segmentation

- In NLP, text is segmented into units of various granularities, which include:
 - Chapters and sections;
 - Paragraphs;
 - Sentences;
 - Clauses;
 - Phrases;
 - Words;
 - Morphemes (stems, suffixes, prefixes).

Text segmentation

- Text segmentation is not straightforward in many cases:
 - For languages like Chinese, Japanese, Tibetan, Thai, there are no spaces between words;
 - For languages like Thai and Tibetan, the delimiters between sentences, clauses or phrases are ambiguous, which makes it hard to segment sentences;
 - Even for English, sentence segmentation is not a trivial task, because the full stop mark (.) is also used for abbreviations, decimals, etc., which may or may not terminate a sentence.

Thai

โลกเราเป็นอะไรหนอในช่วงนี้ ฝั่งหนึ่งของโลกมีอากาศ
อันแปรปรวนวิปริต หนาวเหน็บอย่างไม่เคยเกิดขึ้นมาก่อน
และยังเกิดแผ่นดินพิโรธโกรธคร่าชีวิตคนไปเป็นเรือนแสน
ส่วนบ้านเรานั้นในปีที่ผ่านมาแทบไม่มีฤดูหนาวให้ชื่นใจกันเลย
อากาศกลับร้อน แดดมีทั้งฝนหลงฤดูในช่วงนี้อีกต่างหาก
ทุกคนพูดว่า เป็นเพราะภาวะโลกร้อนนั่นเองที่ทำให้ทุกอย่าง
ดูไม่เหมือนเดิม ประเทศที่มีอากาศหนาวก็หนาวสุดขั้ว ประเทศ

Spaces are not reliable boundaries between sentences.

Chinese

西游记 4 真假猴王

师徒四人继续西行。有一天，他们来到一个地方，前面是望不到边的水面，唐僧发愁 (chóu) 道：“这么大的水，怎么过去呢？”

四个人正不知道怎么办，忽然看见远处好像有一个人在河边，于是就走过去，想问一问。

走近了一看，那不是一个人，而是一块石头，石头上写着三个大字“通天河”，旁边还有一行小字——“河宽 (kuān) 八百里，自古少人行”，意思是这条河有八百里宽，很少有人能通过。

There are not spaces between words.

English sentence segmentation

- Dot marks (.) are ambiguous:
 - Full stop: *This is an apple.*
 - Decimal: *235.6*
 - Abbreviations: *U.S. Ph.D. etc.*
 - A dot mark can take multiple roles: *He comes from U.S.*
- To segment English text into sentences, we need to determine whether a dot mark is an end of sentence or not.
- It can be solved as a classification problem.

English sentence segmentation

— as a classification task

He comes from U.S. She comes from Australia.

↑ ↑ ↑

No Yes Yes

He comes from U.S. with his friends.

↑ ↑ ↑

No No Yes

Chinese word segmentation

(a)	<p>下雨天留客天留我不留</p> <p>下雨、天留客。天留、我不留！</p> <p>下雨天、留客天。留我不？留！</p>	<p>Unpunctuated Chinese sentence</p> <p><i>It is raining, the god would like the guest to stay. Although the god wants you to stay, I do not!</i></p> <p><i>The rainy day, the staying day. Would you like me to stay? Sure!</i></p>
(b)	<p>我喜欢新西兰花</p> <p>我 喜欢 新西兰 花</p> <p>我 喜欢 新 西兰花</p>	<p>Unsegmented Chinese sentence</p> <p><i>I like New Zealand flowers</i></p> <p><i>I like fresh broccoli</i></p>

<http://what-when-how.com/how-to-build-a-digital-library/word-segmentation-and-sorting-digital-library/>

Chinese word segmentation may results in different meanings.

Chinese word segmentation

— as a character tagging task

S	S	B	E	B	M	E	S
我	有	一	台	计	算	机	。
(I)	(have)	(a)		(computer)			(.)

Wang & Xu, Convolutional Neural Network with Word Embeddings for Chinese Word Segmentation, IJCNLP 2017

Tags:

- **S**: single character word
- **B**: beginning character of a word
- **M**: middle character of a word
- **E**: end character of a word

English word segmentation - Tokenization

— A example of Stanford Tokenizer

Input

Another **ex-Golden Stater**, Paul Stankowski from **Oxnard**, is contending for a berth on the **U.S.** Ryder Cup team after winning his first PGA Tour event last year and staying within three strokes of the lead through three rounds of last **month's U.S. Open**. **H.J.** Heinz Company said it completed the sale of its **Ore-Ida** frozen-food business catering to the service industry to McCain Foods Ltd. for about **\$500** million. **It's** the first group action of its kind in Britain and one of only a handful of lawsuits against tobacco companies outside the **U.S.**

Note: **Text in red:** change, **text in blue:** Keep

English word segmentation - Tokenization

— A example of Stanford Tokenizer

Output

Another **ex-Golden Stater** , Paul Stankowski from **Oxnard** , is contending for a berth on the **U.S.** Ryder Cup team after winning his first PGA Tour event last year and staying within three strokes of the lead through three rounds of last **month's U.S. Open** . **H.J.** Heinz Company said it completed the sale of its **Ore-Ida** frozen-food business catering to the service industry to McCain Foods Ltd. for about **\$ 500** million . **It's** the first group action of its kind in Britain and one of only a handful of lawsuits against tobacco companies outside the **U.S.** .

Note: **Text in red:** change, **text in blue:** Keep

Morphological analysis

- To break word down into component morphemes and build a structured representation
- A morpheme is the minimal meaning-bearing unit in a language.
 - **Stem**: the morpheme that forms the central meaning unit in a word
 - **Affix**: prefix, suffix, infix, circumfix
 - **Prefix**: e.g., possible → **im**possible
 - **Suffix**: e.g., walk → walking**ing**
 - **Infix**: e.g., hingi → **hum**ingi (Tagalog)
 - **Circumfix**: e.g., sagen → **ge**sagt**t** (German)

a slide from UW LING 570 by Fei Xia

Two slightly different tasks

- Stemming:
 - Ex: writing \rightarrow writ + ing
- Lemmatization:
 - Ex1: writing \rightarrow write +V +Prog
 - Ex2: books \rightarrow book +N +Pl
 - Ex3: writes \rightarrow write +V +3Per +Sg

Ambiguity in morphology

- flies \rightarrow fly +N +PL
- flies \rightarrow fly +V +3rd +Sg

a slide from UW LING 570 by Fei Xia

Language variation

- Analytic languages: e.g., Chinese; English as a language with analytic tendency.
- Synthetic flexive languages: e.g., Russian
- Synthetic agglutinate languages: e.g., Turkish

Ways to combine morphemes to form words


- Inflection: stem + gram. morpheme → same class
 - Ex: help + ed → helped
- Derivation: Derivation: stem + gram. morpheme → different class
 - Ex: civil + -zation → civilization
- Compounding: multiple stems
 - Ex: cabdriver, doghouse
- Cliticization: stem + clitic
 - Ex: they'll, she's (*I don't know who she is)

a slide from UW LING 570 by Fei Xia

Finite state transducers (FSTs)

- Finite State Transducers are an extension to Finite State Machines, where an output symbol will be given for each input symbol.
- FSTs are commonly used tools for morphological analysis.
- A FST can be used in a inverse direction with the input and the output swapped.

Finite state transducers (FSTs)



input	output
0	0
11	01
110	010
1001	0011
1100	0100
1111	0101
10010	00110

English morphology

- Affixes: prefixes, suffixes; no infixes, no circumfixes.
- Inflectional:
 - Noun: -s
 - Verbs: -s, -ing, -ed, -ed
 - Adjectives: -er, -est
- Derivational:
 - Ex: $V + \text{suf} \rightarrow N$
computerize + -ation \rightarrow computerization
kill + er \rightarrow killer
- Compound: pickup, database, heartbroken, etc.
- Cliticization: 'm, 've, 're, etc.


a slide from UW LING 570 by Fei Xia

Three components

- Lexicon: the list of stems and affixes, with associated features.
 - Ex1: book: N
 - Ex2: -s: +PL
- Morphotactics:
 - Ex: +PL follows a noun
- Orthographic rules (spelling rules): to handle exceptions that can be dealt with by rules.
 - Ex3: $\epsilon \rightarrow e / x^{\wedge} _ s\#$

a slide from UW LING 570 by Fei Xia

Rewrite rules



An example

Task: foxes → fox +N +PL

Surface: foxes



Orthographic rules

Intermediate: fox ^s




Lexicon + morphotactics

Lexical: fox +N +pl

a slide from UW LING 570 by Fei Xia

An FST




cat +N +PL \rightarrow cat \wedge s #

cat +N +Sg \rightarrow cat #

a slide from UW LING 570 by Fei Xia


Expanding FST



fox +N + Pl \rightarrow fox \wedge s #
 cat +N + Pl \rightarrow cat \wedge s #
 goose +N +Sg \rightarrow goose #
 goose +N +Pl \rightarrow geese #

a slide from UW LING 570 by Fei Xia

Representing orthographic rules as FSTs



$\epsilon \rightarrow e / (s|x|z) \hat{_} s \#$


Input: $\dots(s|x|z) \hat{s} \#$ immediate level

Output: $\dots(s|x|z)es \#$ surface level

To reject (fox \hat{s} , foxs)

a slide from UW LING 570 by Fei Xia

Representing orthographic rules as FSTs



(fox, fox): q0, q0, q0, q1

(fox#, fox#): q0, q0, q0, q1, q0

(fox^z#, foxz#): q0, q0, q0, q1, q2, q1, q0

(fox^s#, foxes#): q0, q0, q0, q1, q2, q3, q4, q0

(fox^s, foxs): q0, q0, q0, q1, q2, q5

a slide from UW LING 570 by Fei Xia

Further reading on morphological analysis

- Fei Xia, slides on morphological analysis

https://www.powershow.com/viewfl/6a39a-ZDc1Z/Morphological_analysis_powerpoint_ppt_presentation

- Mans Hulden (2011), Morphological analysis with FSTs

<https://fomafst.github.io/morptut.html>

Content

- 1 About the course
- 2 Research questions and NLP tasks
- 3 Grammars and Automata
- 4 Text segmentation and morphology analysis
- 5 Word frequency and collocations**

Top 5000 words in American English


Rank	Word	Part of speech	Frequency	Dispersion
1	the	a	22038615	0.98
2	be	v	12545825	0.97
3	and	c	10741073	0.99
4	of	i	10343885	0.97
5	a	a	10144200	0.98
6	in	i	6996437	0.98
7	to	t	6332195	0.98
8	have	v	4303955	0.97
9	to	i	3856916	0.99
10	it	p	3872477	0.96

Rank	Word	Part of speech	Frequency	Dispersion
1	the	a	22038615	0.98
2	be	v	12545825	0.97
3	and	c	10741073	0.99
4	of	i	10343885	0.97
5	a	a	10144200	0.98
6	in	i	6996437	0.98
7	to	t	6332195	0.98
8	have	v	4303955	0.97
9	to	i	3856916	0.99
10	it	p	3872477	0.96

Statics from Corpus of the Contemporary American English

<http://www.wordfrequency.info/>

Top 5000 words in American English




Zipf's Law

The frequency of any word is inversely proportional to its rank in the frequency table:

$$p(w_r) \propto \frac{1}{r}$$

Zipf's law



A plot of the rank versus frequency for the first 10 million words in 30 Wikipedias (dumps from October 2015) in a log-log scale.

(By SergioJimenez - Own work, CC BY-SA 4.0, from Wikipedia)

Collocation or multi-word expression (MWE)

- A COLLOCATION is an expression consisting of two or more words that correspond to some conventional way of saying things.
- The words together can mean more than their sum of parts
 - The Times of India, disk drive
 - hot dog, mother in law

Collocation or multi-word expression (MWE)

- Examples of collocations
 - noun phrases like *strong tea* and *weapons of mass destruction*
 - phrasal verbs like to *make up*, and other phrases like the *rich and powerful*.
- Valid or invalid?
 - *a stiff breeze* but not a *stiff wind* (while either a *strong breeze* or a *strong wind* is okay).
 - *broad daylight* (but not bright daylight or narrow darkness).

Manning & Schütze, Fundamentals of Statistical Natural Language Processing, 1999

Criteria for collocations (or MWE)

- Typical criteria for collocations:
 - non-compositionality
 - non-substitutability
 - non-modifiability.
- Collocations usually cannot be translated into other languages word by word.
- A phrase can be a collocation even if it is not consecutive (as in the example *knock ... door*).

Manning & Schütze, Fundamentals of Statistical Natural Language Processing, 1999

Non-Compositionality

- A phrase is compositional if the meaning can be predicted from the meaning of the parts.
 - E.g. new companies
- A phrase is non-compositional if the meaning cannot be predicted from the meaning of the parts
 - E.g. *hot dog*

Manning & Schütze, Fundamentals of Statistical Natural Language Processing, 1999

Non-Compositionality

- Collocations are not necessarily fully compositional in that there is usually an element of meaning added to the combination.
 - E.g. *strong tea*
- Idioms are the most extreme examples of non-compositionality
 - E.g. *to hear it through the grapevine*

Manning & Schütze, Fundamentals of Statistical Natural Language Processing, 1999

Non-Substitutability

- We cannot substitute near-synonyms for the components of a collocation.
- For example
 - We can't say *yellow wine* instead of *white wine* even though *yellow* is as good a description of the color of *white* wine as white is (it is kind of a yellowish white).

Manning & Schütze, Fundamentals of Statistical Natural Language Processing, 1999

Non-Substitutability

- Many collocations cannot be freely modified with additional lexical material or through grammatical transformations (Non-modifiability).
 - E.g. *white wine*, but not *whiter wine*
 - E.g. *mother in law*, but not *mother in laws*

Manning & Schütze, Fundamentals of Statistical Natural Language Processing, 1999

Metrics for Collocation or MWE Extraction

- Frequency
- Mean and Variance of Distances between Words
- Hypothesis Testing
 - t -test
 - χ^2 test
 - likelihood ratio test
- Mutual Information
- Left and Right Context Entropy
- C-Value

Further reading on collocation and MWE

- Manning & Schütze, Fundamentals of Statistical Natural Language Processing, 1999, Chapter 3 (A general introduction to collocation)
- Katerina T. Frantzi, Sophia Ananiadou, Junichi Tsujii, The C-value / NC-value Method of Automatic Recognition for Multi-word Terms, ECDL 1998: Research and Advanced Technology for Digital Libraries pp 585-604 (proposed the C-value metric)
- Zhiyong Luo, Rou Song, An integrated method for Chinese unknown word extraction, SIGHAN 2004. Barcelona, Spain. (proposed the context entropy method)

Content

- 1 About the course
- 2 Research questions and NLP tasks
- 3 Grammars and Automata
- 4 Text segmentation and morphology analysis
- 5 Word frequency and collocations